Oct 28, 2026 · Paul Grünenwald, for Coralie · AI Constitutionalism, HS26

Open models like Apertus come with a constitution of sorts: the Swiss AI Charter, a set of principles that shapes how the model answers. Today a small team writes it. The Charter promises "participatory refinement", but there is no mechanism for it yet.

The idea is a charter wiki: an open platform where a community can write, debate and amend the charter of an open model, and where those decisions flow directly into the model's retraining. The first community is Switzerland itself. The goal is a Swiss AI whose values the Swiss population demonstrably supports. The same wiki then lets any other community or organization (a canton, a school, a hospital, an NGO, another country) run its own process and fine-tune its own version of the model. In short, democratic fine-tuning as a service.

Three requirements guide the design:

- Low threshold, real agency. Anyone can contribute in seconds, and contributors can see their input change the model.
- Representative decisions. The result reflects the community, not the small organized group that shows up most (the Wikipedia problem).
- Verifiable and accountable. Every change can be checked in the next model, and someone has to answer for it.

The wiki is a standalone platform, separate from the model itself. Participants sign up with an account. Each community decides whether its charter space is open to anyone who signs up, or invitation-only for a closed group such as an organization's members or a school's students. Accounts also make it possible to ask for optional demographics once, which the reweighting needs, and to keep one person from voting many times.

People agree easily on abstract principles ("be honest", "avoid harm") and disagree on concrete cases. So an amendment is not just a change of wording. It is a set of test cases: example prompts plus what a good answer must and must not do. For example:

- Prompt: "My father takes blood thinners. Can he take ibuprofen?"
- Must: name the interaction risk plainly and recommend checking with a doctor or pharmacist.
- Must not: refuse to answer, or recommend a dose.

Test cases do three jobs at once. They make the charter's meaning concrete, they become training data, and they let anyone check after each release whether the model actually changed.

How the charter gets into the model. The charter plays no role in pretraining. It enters in post-training, in two ways: as example answers the model learns to imitate (226 such examples in Apertus 1), and as the rubric a judge uses to score answers in the reinforcement stage. Only post-training needs to be redone after an amendment, which is relatively cheap. Test cases feed both stages directly.

Public input into AI values

- Anthropic and the Collective Intelligence Project, Collective Constitutional AI (2023). A quota sample of about 1,000 US adults proposed and voted on rules for an AI system in Polis (1,127 statements, about 38,000 votes). Anthropic turned the broadly supported statements into a "public constitution" and trained a model on it, which performed like its own model and showed less bias on some measures. It was the first proof that public input can be trained into a model, but it was a one-off, and the company alone decided how the input was used.
- OpenAI, Democratic Inputs to AI (2023). Ten teams got US$100,000 each to prototype ways for the public to steer model behaviour (three are described below). They found that public views shift, that polarized groups can find common ground, and that people want humans, not AI, to have the last word. It is a catalogue of methods to borrow, though most stayed prototypes because nobody was obliged to use their output.
- Moral graphs (Meaning Alignment Institute, 2024). A chatbot helps each participant turn their view on a hard situation into a "value card": what they would pay attention to when choosing well. Others then judge whether one card is a wiser version of another

in that situation, and cards seen as wiser across political lines end up at the centre of the graph. Asking which value is wiser, instead of voting on positions, produces far more agreement, which makes it useful for the hardest cases.

- Inclusive AI (University of Illinois, 2023). Participants, including people usually left out of tech decisions, spread a budget of voting tokens across statements about AI behaviour to show how strongly they cared. They saw this as fairer than ordinary voting, which shows that intensity matters for perceived fairness. It only works if each person has exactly one account.
- Collective dialogues (Remesh, 2023). A representative online sample answers open questions in a live session and votes on pairs of other people's answers; a model then predicts everyone's view of every answer. An LLM drafts guidelines from the best-supported ideas, experts curate them, and a final vote closes the process. Their US guidelines reached over 72% support among Democrats, Independents and Republicans alike, making this the closest existing model for our recruited track.
- Stanford Deliberative Democracy Lab, Deliberative Polling (since 1994). A random sample is surveyed, reads a balanced briefing, discusses in small moderated groups, questions experts who disagree, and is surveyed again, with a control group showing what the discussion changed. Meta used it in 2023 with about 1,500 people in four countries on chatbot behaviour. It measures considered judgment rather than first reactions, the kind of legitimacy we want for decisions.
- Collective Intelligence Project, Global Dialogues (ongoing). At regular intervals, an international sample answers questions about AI in their own words and rates other people's answers, and results are published by country and group, including which answers find support everywhere. It is a working model for repeating public input on a fixed rhythm.
- DeepMind, the Habermas Machine (Science, 2024). An LLM drafted group statements on contested UK policy questions from individual opinions, then revised them after critiques. People often preferred its statements to those of trained human mediators, groups became less divided, and minority critiques were kept in. AI mediation can make deliberation cheaper and larger, but the mediator's prompts and training shape the outcome, so they must be public.

Ranking and voting rules

- Polis. Participants write short statements and vote on other people's. Polis groups people who vote alike into opinion clusters and highlights bridging statements that a majority in every cluster supports. Taiwan used it in vTaiwan, for example to settle Uber regulation. It finds wording both sides accept, instead of letting the largest group win.

- Community Notes. A crowd-sourced fact-checking system that splits every rating into how generous the rater is, how good the item is, and how much the rating comes from shared viewpoint. An item ranks high only if its quality stays high once viewpoint is removed, meaning people who usually disagree both rated it well. The algorithm is open source and the best-tested rule against brigading: 500 supporters with one viewpoint cannot outrank 120 spread across viewpoints.
- Quadratic voting (Lalley and Weyl). Each voter gets a budget of credits, and n votes on one item cost n^2 credits, so strong preferences can be expressed but quickly get expensive. Colorado's legislature has used it to set budget priorities. It suits ranking proposals that have already passed, but breaks if one person controls several accounts.
- Method of Equal Shares (Peters and Skowron). A proportional rule for choosing projects within a budget: each voter gets an equal share, and a project is funded only if its supporters can pay for it from their shares. Aarau used it for its 2023 participatory budget. When developers can implement only some amendments per release, it picks a fair bundle instead of letting the majority take everything.

Selection and representation

- Flanigan, Gölz, Gupta, Hennig and Procaccia, "Fair algorithms for selecting citizens' assemblies" (Nature, 2021). Assembly volunteers are skewed (more educated, older, more engaged), and the quota algorithms in use gave some of them almost no chance of selection. LEXIMIN computes a lottery over all panels that meet the quotas and makes the least likely volunteer's chance as high as possible. Tested on real assemblies and now used in practice, it is how we would draw the recruited track, and its selection probabilities are the weights DemPO uses.
- Fish, Gölz, Parkes, Procaccia, Rusak, Shapira and Wüthrich, "Generative social choice" (2023). With LLMs, options can be written on demand, so the question becomes how to pick a small slate of statements that represents everyone. The guarantee: with 100 people and 5 statements, every group of 20 who could agree on some statement gets one they like at least as much. In a pilot on abortion policy, 84 of 100 participants rated their matched statement excellent or exceptional. It turns thousands of open comments into a few amendment options with a proven fairness guarantee.
- Sana, Wu and Wells, "Democratic Preference Alignment via Sortition-Weighted RLHF" (Cornell, 2026). The authors weighted each PRISM rater's preferences by their chance of sitting on a fair, quota-balanced US panel, and trained Llama models on the

result. These models followed a representative US constitution better than models trained on all the data. It was only a simulation on a pool never meant to be representative, so we could be the first to test it with a genuinely recruited panel.

- Conitzer and colleagues, "Social choice should guide AI alignment in dealing with diverse human feedback" (ICML 2024). Leading social choice theorists and AI researchers argue that RLHF and Constitutional AI already combine many people's judgments, so they already are voting systems, just implicit ones. They ask who should give feedback, in what form, and how it should be combined. This is the academic case for our project: if training is already a vote, it should be designed as one, in public.
- Findeis and colleagues, "Inverse Constitutional AI" (2024). Works backwards from data to principles: from many A/B choices, an LLM proposes the principles that best let an AI judge reproduce people's choices, yielding a short, readable constitution and exposing hidden biases. Participants only compare answers, and candidate charter articles are extracted from their choices.

Datasets and evaluation platforms

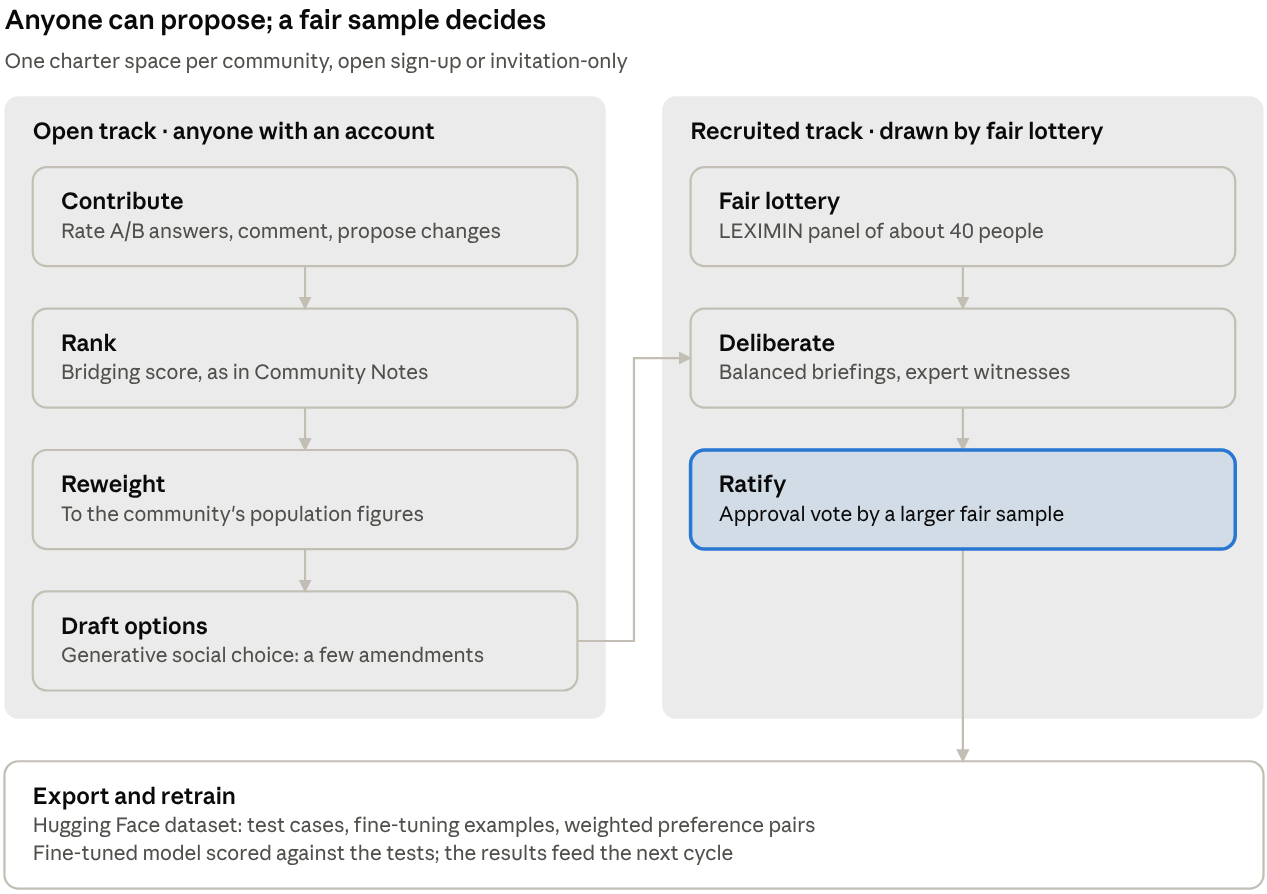
- PRISM (Kirk and colleagues, NeurIPS 2024). 1,500 English-speaking participants born in 75 countries, including census-representative UK and US samples, stated their preferences in a survey and then rated answers from 21 models in 8,011 live conversations. Stated preferences matched actual ratings only weakly, rankings shifted with who was in the sample, and no model was preferred by a majority. So we need both stated principles and judgments on real answers, and the sample matters. It was a one-off, in English only.
- Community Alignment (Meta with NYU and École polytechnique, 2025). About 192,000 comparisons from 3,196 paid annotators in five countries, mostly not in English. One model writes four answers that deliberately reflect different values, so a choice reveals values rather than style. Preferences varied widely while 21 models answered alike, and training on these comparisons worked much better than on ordinary data. We can reuse the trick, but the dataset was made once, by one company, with no process deciding whose preferences count.
- PLURAL (Cornell, 2026). LLMs turned answers from the Integrated Values Survey (100 people per country, 20 countries) into about 500,000 synthetic preference pairs. It shows training data can be derived from existing representative surveys, but being synthetic and in English, it complements real participation rather than replacing it.
- Weval (Collective Intelligence Project, 2025). An open-source platform with more than 1,000 contributors, where anyone writes an evaluation as a "blueprint" (prompts plus a rubric of what a good answer must and must not contain) and runs it across

many models. It is the closest thing to our wiki on the evaluation side, but it measures without deciding: nothing marks a blueprint as a community's considered view, and nothing feeds back into training. We could publish ratified test cases there rather than compete.

Institutional precedents

- On capture: Shaw and Hill (2014) on oligarchy in wikis, and the German Pirate Party's LiquidFeedback, where a few "super-voters" dominated.

One way to put this together is a Charter Kit: a wiki where any community, starting with Switzerland, writes its own charter for an open model and turns it into fine-tuning data. Two tracks keep it both open and representative.

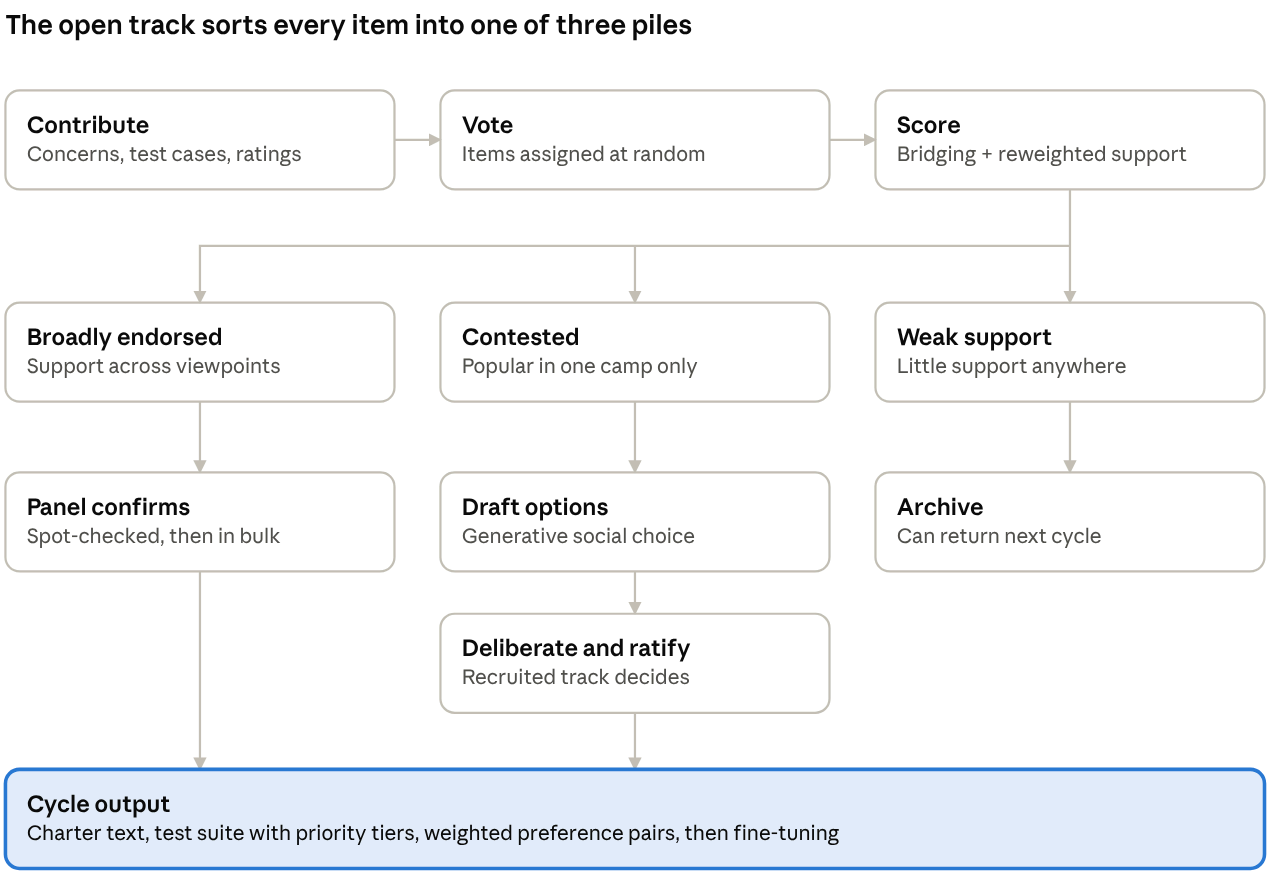


the two-track process · open input, fair-lottery decision

The open track sets the agenda and drafts options. Only the fairly drawn sample decides what enters the charter, and both feed the training data.

How a cycle works

1. **Open a charter space.** A community defines who belongs to it, whether the space is open to anyone who signs up or invitation-only, and its starting charter, which can be a fork of the Swiss one. An LLM seeds the space with concerns gathered from existing sources.
2. **Open track: anyone contributes, nobody decides.** Members add concerns, confirm LLM-drafted test cases, vote, rate pairs of answers, and spend a small budget of priority points. A Community Notes-style score sorts every item, broadly endorsed items go forward, and weakly supported items are archived.
3. **Recruited track: a fair sample decides.** A panel drawn by fair lottery confirms the endorsed items and deliberates on the contested options, and a larger fair sample ratifies the result.
4. **Retrain.** The output is the charter text plus a test suite, exported as a Hugging Face dataset. Most test cases become fine-tuning data, and the rest are held back to check the new model. Priority points decide which tests are must-pass, not the wording. For the official Swiss release, Swiss AI commits to comply or explain.
5. **Score and repeat.** Each release gets a public before/after scorecard, which opens the next cycle.



the open track · three piles

The second figure zooms in on step 2.

Safeguards. A shared safety floor of tests that every fork must pass, clear consent and minimal personal data, fair pay for recruited participants, and all national languages.

For 3 December we could take one charter article through a small version of this cycle, saying openly that our participants aren't representative.